

Did Generative AI Trigger a Structural Break? ARIMA Analysis of XLK ETF

1 Introduction

Technology stocks have accounted for 32% of global equity returns and 40% of total U.S. equity market returns since 2010 (Goldman Sachs, 2024), highlighting the dominant role of the technology sector in modern financial markets. Large-cap firms, such as Apple and Microsoft, represent a substantial share of total market capitalization; consequently, their performance increasingly shapes the trajectory of major stock indices, including the S&P 500 and the NASDAQ Composite. Because technology companies are largely valued based on expectations of future innovation and scalability, breakthroughs can trigger significant revisions in investor sentiment, often leading to substantial movements in stock prices.

The introduction of new technology significantly impacts a firm's business model. For technology-centric companies in particular, launching a novel product often increases the company's valuation (McKinsey, 2016). As several tech giants have launched flagship AI products, such as Google's Gemini and Anthropic's Claude, their market valuations and stock prices have shifted accordingly. Although OpenAI is not publicly traded, the "hard launch" of ChatGPT in November 2022 significantly increased OpenAI's private valuation and attracted substantial investment, with its estimated worth skyrocketing from roughly \$20 billion to \$85.2 billion by 2026 (Barrientos, 2026).

The adoption of AI applications has reshaped market narratives surrounding technological growth and productivity. Beyond the firms directly launching AI models, companies perceived as key beneficiaries of AI infrastructure, most notably NVIDIA Corporation, have also experienced significant revaluations. This shift suggests that innovation shocks may influence not only short-term prices but also the broader structural dynamics of the technology sector.

Understanding whether such a technological breakthrough corresponds to a lasting change in market behavior is critical for investors and policymakers. Persistent shifts in return patterns can affect risk assessment, portfolio allocation, and forecasting models. Financial return series

are time-dependent and often exhibit changes in volatility and growth dynamics. If the emergence of generative AI has altered investor expectations in a sustained way, this effect may manifest as a structural break in the return-generating process.

This paper investigates whether the emergence of generative AI signaled a fundamental change in the behavior of technology-sector stock returns. We evaluate this change by comparing stock performance before and after the launch of ChatGPT. Specifically, we analyze the Technology Select Sector SPDR Fund (XLK) using ARIMA models to examine whether the release of ChatGPT corresponds to a statistically significant shift in XLK price dynamics from 2017 to 2024.

We chose the XLK because it's a comprehensive fund that captures an effective representation of the technology sector of the S&P 500 Index. The XLK has a 4-star in the Morningstar rating by March 2026, indicating an above-average performance that the XLK is in the top 22.5% to 32.5% of technology category. The XLK ETF provides precise exposure to companies from technology hardware, storage and peripherals; software; communications equipment; semiconductors and semiconductor equipment; IT services; and electronic equipment, instruments and components industries (State Street). This feature allows investors to take strategic or tactical positions at a more targeted level than traditional style based investing. The companies in the XLK are selected from the S&P 500 based on the Global Industry Classification Standard (GICS) where companies are assigned to a single classification based on revenue.

2 Data

The XLK stock price we are looking at is collected by S&P Dow Jones Indices and State Street Global Advisory. S&P Dow Jones Indices is in charge of creating the company benchmark on entering the technology company stock pool. The State Street Global Advisors selected particular stocks into the combination of the XLK stock. The data we are using is collected from the investing.com website, a financial data platform that provides public real-time market data directly from stock exchanges such as the New York Stock Exchange and NASDAQ.

2.1 Data Overview

Using the dataset from the investing website, we derive the following variables for the analysis:

- **Price** (in USD): The final price at which the ETF traded when the stock exchange closed (typically 4:00 PM EST).
- **Open** (in USD): The price of the very first transaction when the market opened (9:30 AM EST).

- **High & Low** (in USD): The maximum and minimum prices reached during the trading session.
- **Vol.:** The total number of shares traded during the day (e.g., 15.90M means 15.9 million shares).
- **Change:** The percentage difference between the current day’s closing price and the previous day’s closing price.

The dataset comprises daily observations spanning from January 3, 2017, to March 28, 2026. This time frame was specifically selected to ensure that the launch of ChatGPT (November 30, 2022) serves as a central pivot point, providing a balanced sample size for comparing pre- and post-innovation market dynamics

2.2 Data Cleaning

In the original dataset downloaded from the investing website, we observed outliers that are doubled the price of the remaining data points. We also found that for the weekends, prices are carried over from the Friday of the same week. These data points correspond to weekend days, and they have no data for the variable, Vol.. We decided to remove these points because they are not consistent with the general pattern, which they doubled the Friday price. Lastly, we split the dataset into two groups: `before_data` and `after_data`, using November 30, 2022 (the release date of ChatGPT) as the cutoff.

Table 1: Summary Statistics of Stock Price

Table 1: Summary Statistics of Stock Price

mean_price	sd_price	min_price	max_price
69.91683	33.95005	24.4	152.07

Table 1 shows the mean, standard deviation, minimum and maximum of the stock price between 2017/1/3 to 2026/3/27.

Section 3. Analytic Methods

Figure 1 shows the stock price over time. The red dashed line indicates the time when ChatGPT was released. Overall, the stock price shows a general upward trend during the period shown.

Figure 2 indicates that the time series is not white noise, as the ACFs are out of the boundary and decay very slowly, which looks like a unit root process. Thus, we need to transform the




Figure 1




Figure 2

data to make it stationary before model fitting. For stock data, taking the natural log is a common way to transform the data. The first difference of the log price series, known as the log return, is then used to remove the trend and produce a stationary series suitable for model estimation.




Figure 3

Figure 3 shows the log return over time after transformation. The log return fluctuates around zero, and the fluctuations are small for most of the time. However, there are a few large spikes. The biggest fluctuations occur around 2020 when the pandemic began. This suggests that the pandemic caused substantial uncertainty and sharp movements in the market.

Figure 3 suggests the mean is around zero, which suggests log return time series might be stationary. But Figure 4 shows the ACF of log return, where several spikes that are not within the bound. Therefore, an Augmented Dickey-Fuller (ADF) test was conducted to formally determine whether the series is stationary.

The ADF test has a p value of 0.01, suggesting log return is stationary.

3 Method

3.1 Assess difference between before AI and after AI

ACF of log return




Figure 4

PACF of log return




Figure 5

Table 2: Summary Statistics of Stock Price

period	Mean	SD
After AI	-0.00078	0.01456
Before AI	-0.00069	0.01641

According to table Table 2, the mean log return before and after the appearance of ChatGPT is both close to zero. To test whether there is a difference on price before and after ChatGPT appearance, a regression model with a dummy variable D_t as the predictor was fitted.

$$Price = \beta_0 + \beta_1 + \beta_2 D_t$$

where $D_t = 0$ is before the appearance of ChatGPT, and $D_t = 1$ is after the appearance of ChatGPT. A two sample mean test can be used to determine whether D_t is significant. The null hypothesis is that the mean price in the two periods are equal, while the alternative hypothesis is that they are different.

Table 3: Model Coefficients

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	14.360	0.486	29.554	0.000
time_index	0.047	0.001	85.342	0.000
AI_dummy	2.039	0.773	2.638	0.008

The p value of the dummy variable is 0.008, which is greater than 0.001. Thus, we fail to reject the null hypothesis, and conclude that there is no difference on the price values before and after the appearance of AI. From Figure 6, we do not see a significant difference between the regression line of the two time periods.

3.2 Fit ARIMA Model

Using the ACF and PACF plots, we can determine the appropriate order of the ARIMA model. According to Figure 4, the ACF gradually tails off, suggesting the presence of autoregressive behavior. According to Figure 5, there are strong spikes at lags 2, 4, 7, and 8, and several other significant lags. Based on these patterns, ARIMA(1,1,1), ARIMA(2,1,2), ARIMA(4,1,4), ARIMA(7,1,7), and ARIMA(8,1,8) models were fitted to the price data to capture the time series.









Figure 6







Model Diagnosis

The p values for Ljung-Box Statistic plots indicate that ARIMA(1,1,1), ARIMA(2,1,2), and ARIMA (4,1,4) are not good fits, because p values are significant for all lags. ACF plots for ARIMA(7,1,7) and ARIMA (8,1,8) looks stationary, but QQ plots show the points are deviated from the reference line, meaning the normality condition is not met. To determine which model performs the best, we compare the model AIC, AICc, and BIC values.

Model Choice

Table 4: Comparison of Model Performance

	AIC	AICc	BIC
ARIMA(1,1,1)	-5.477657	-5.477655	-5.470223
ARIMA(2,1,2)	-5.492952	-5.492945	-5.480561
ARIMA(4,1,4)	-5.503769	-5.503742	-5.481465
ARIMA(7,1,7)	-5.503271	-5.503192	-5.466098
ARIMA(8,1,8)	-5.501304	-5.501202	-5.459175

The ARIMA(7,1,7) has the smallest AICc value among all the models. Therefore, we pick ARIMA(7,1,7) as our best model.

Forecasting

Benchmark Methods

We are interested in how significant the data after November 30, 2022, when ChatGPT got launched, is to forecasting future stock market prices. We will be using data frames of different time sets: January 3, 2017 to March 27, 2026, January 3, 2017 to November 30, 2022, and November 30, 2022 to March 27, 2026. The forecast models we will be using three of the “benchmark” methods: Mean method, Naive method and Seasonal naive method. Mean method forecasts all future values as equal to the average (or mean) of the training set. Naive method set all forecast values to be the value of the last observation of the historical data. Seasonal naive method set each forecast to be equal to the last observed value from the same season. We tested if there is a weekly seasonality. Drift method and decomposition method, which both looks at the amount of change over time or the trend, seem unnecessary with log return data because we have already applied differencing.

First, we aimed to examine whether AI has had an unexpected impact on the stock market and contributed to its fluctuations. We would like to use data from January 3, 2017 to November 30, 2022 to forecast data after November 30, 2022.

Table 5: Forecast Errors

Table 5: Forecast Errors

Method	MAE	RMSE
Mean	0.01047	0.01455
Naive	0.04840	0.05036
Seasonal	0.02094	0.02603



According to Table 5 MAE (Mean Absolute Error) and RMSE (Root Mean Squared Error), which measure the difference between actual and predicted values, were small for the mean method but also all these benchmark methods. It appears that the data prior to November 30, 2022 can effectively predict the data following the cutoff point. Furthermore, this suggests that AI have not yet contribute a significant impact in the stock market. As a result, we will use the entire time frame, allocating 80% for training and 20% for testing, to build a forecast model in predict prices in future 2026.

Table 6: Forecast Errors-full model

Table 6: Forecast Errors-full model

Method	MAE	RMSE	MASE	MASSE
Mean	0.01125	0.01628	0.71590	0.68334
Naive	0.01118	0.01632	0.71095	0.68521
Seasonal	0.01351	0.01916	0.85955	0.80430

Table 6 shows the naive method has the lowest MAE and mean method is has the lowest RMSE. MAPE (Mean Absolute Percentage Error) is a percentage-based error metric that scales forecast errors relative to the actual values. We did not include MAPE here for comparison, because our log return data moves around 0, where some small values here make MAPE enormous and not as a good parameter to look at. MASE (Mean Absolute Scaled Error) and MASSE (Root Mean Squared Scaled Error) are scaled error metrics, and they works by scaling errors based on the training MAE from a simple forecast method. This makes them stable because the denominator is a fixed constant derived from the training data, providing a more informative alternative when MAPE fails to give meaningful results. Lowest MASE and MASSE suggest the Mean method is the best, so we will use mean method to do forecasting.



Next, we construct a prediction interval for the mean method forecast model. While the histogram suggests the residuals are approximately normally distributed, the QQ-plot indicates that the log returns deviate from normality, violating classic assumptions. This implies that the resulting prediction interval might not be reliable.




Figure 7

ARIMA forecasts

Using $\text{ARIMA}(7,1,7)$, we want to forecast the data in the next 30 days starting March 28, 2026.




Figure 8




Figure 9

Summary

This project tests whether the rise of AI has significantly affected the behavior of technology sector stock returns. Using the XLK ETF price data from 2019 to 2026, we transformed the series into log returns to make the data stationary and applied AR(1), MA(1), and ARMA (1,1) time series models. The result suggests that stock return pattern remained relatively stable before and after the emergence of AI, indicating that AI has not yet had a significant impact on the technology stock market price. More broadly, this shows that while AI is often seen as a transformative force, its impact on financial markets may not be immediately reflected in stock return fluctuations.

Limitations

There are some potential limitations for the dataset. One of the limitations is that there's possibility for State Street to rebalance the stock composition over time, which means the dataset may not remain fully consistent throughout the study period. The companies in the XLK ETF are selected based on the GICS. However, the GICS structure is reviewed annually, meaning State Street need to adjust the XLK portfolio at least on a yearly base. This could reduce the reliability of comparing stock performance across different years and weaken the robustness of analysis.

Another limitation is the dependency of the dataset. In Section 3, we found out the time series data points remain highly correlated even after the data transformation. For instance, the QQ plot indicates deviations from normality, suggesting that model assumptions may not hold and prediction intervals may be unreliable. In addition, our model focuses on the short-term impacts after the release of ChatGPT, but the effect of AI may be better reflected in a long term trend rather than a boom. The log transformation may mask the underlying long term trend, because it removes the trend to make the data stationary. The log transformation only focuses on daily changes rather than long-term structural changes. As a results, model based on the log transformation is better at studying short-term market reaction but may be less effective at capturing long-term changes in the market over time.

Also, AR, MA, ARMA are all models that are only a function of the time series itself. They do not include other important factors, such as government policy or economic policy changes. There are many other factors that affects the stock market that we did not incorporate, such as economic indicators and interest rates. Since stock market is influenced by many external variables, excluding these factors may limit the strength and forecasting accuracy of the models.

Future Directions

To account for external factors that might affect the stock market, we could add these as predictors in our model. This can help us better estimate stock prices.

Although we have not been able to conclude AI has significant impact on the behavior of technology sector stock returns, it is still useful to estimate how sensitive technology sector stock price comparing to the broader market indices or other sectors. We might expect the technology sector to show higher volatility and respond more to growth expectations. In future studies, we can fit and compare separate models for different market segments, such as the S&P 500 and individual sectors like healthcare or financials, to better understand these differences.

It is also possible to incorporate the method proposed by Campos-Martins (2024) with our model. After identifying notable development by AI companies, we can adjust our predictors to improve accuracy. For example, we could include an indicator variable that captures major AI events or achievements. This would allow the model to account for changes over time driven by new innovations, instead of relying only on fixed predictors.

Finally, our research should follow a consistent process where we constantly refit the models as the stock market changes over time. Updating the model with new data can help us see whether there are any long-term effects or clear shifts in the data, especially as AI continues to develop and become more widely used.