

Did Generative AI Trigger a Structural Break?

ARIMA Analysis of XLK ETF

1 Introduction

Artificial Intelligence (AI) has been a hot topic in recent years, with its wide applications in scientific discovery, advanced cybersecurity, and personalized AI agents. In the stock market, technology stocks have accounted for 32% of global equity returns and 40% of total U.S. equity market returns since 2010 (Goldman Sachs, 2024), playing a dominant role in modern financial markets. Large-cap firms, such as Apple and Microsoft, represent a substantial share of total market capitalization; consequently, their performance increasingly shapes the trajectory of major stock indices, including the S&P 500 and the NASDAQ Composite (Yahoo Finance). Since Technology companies are largely valued based on expectations of future innovation and scalability, breakthroughs can trigger significant revisions in investor sentiment, leading to movement in stock prices.

Breakthrough in new technology also impacts a firm's business model. For technology-centric companies in particular, launching a novel product often increases company's net valuation (McKinsey, 2016). As several tech giants have launched flagship AI products, such as Gemini from Google and Claude from Anthropic, their market valuations and stock prices have shifted accordingly (Investing.com, 2026). The impact of novel AI product is seemingly significant since it can impact the company's valuation that the company does not even need to be publicly traded. The hard launch of ChatGPT in November 2022 significantly increased OpenAI's, a non-profit artificial intelligence research company, private valuation and attracted substantial investment, with its estimated worth skyrocketing from roughly \$20 billion to \$85.2 billion by 2026 (Barrientos, 2026).

Furthermore, the adoption of AI applications has also reshaped market narratives surrounding technological growth and productivity. Beyond the firms that directly launch AI models, companies perceived as key beneficiaries of AI infrastructure, most notably the GPU manufacturing company NVIDIA Corporation, have also experienced significant revaluations (Morningstar, 2026). This shift suggests that innovation shocks may influence not only short-term prices but also the broader structural dynamics of the overall technology sector.

With the crucial role of technology stock in the general financial market and AI product's impacts on company's evaluation, we want to investigate whether the technological breakthrough impact the technology sector stock market. This report investigates:

- Did the stock return pattern changes after the emergence of AI products?
- What's the predicted stock return based on current stock market trend?

Specifically, we evaluate this change by analyzing the Technology Select Sector SPDR Fund (XLK) using ARIMA models to examine whether the release of ChatGPT corresponds to a statistically significant shift in XLK price dynamics from 2017 to 2024.

Understanding whether such a technological breakthrough corresponds to a lasting change in market behavior is crucial for investors and policymakers. Persistent shifts in return pattern can affect risk assessment, portfolio allocation, and forecasting models in financial market analysis. Financial return series are time-dependent and often exhibit changes in volatility and growth dynamics. If the emergence of artificial intelligence has altered investor expectations, this effect may manifest as a structural break in the return-generating process.

2 Methodology

2.1 Data Overview

This report uses the daily XLK fund return price from Investing.com, a financial data platform that offers real-time data. The XLK stock price is collected by S&P Dow Jones Indices and State Street Global Advisory. S&P Dow Jones Indices is in charge of creating the company benchmark on entering the technology company stock pool. The State Street Global Advisors selected particular stocks into the combination of the XLK stock.

We chose the XLK because it's a comprehensive fund that captures an effective representation of the technology sector of the S&P 500 Index. The XLK has a 4-star in the Morningstar rating by March 2026, indicating an above-average performance that the XLK is in the top 22.5% to 32.5% of technology category. The XLK ETF provides precise exposure to companies from technology hardware, storage and peripherals; software; communications equipment; semiconductors and semiconductor equipment; IT services; and electronic equipment, instruments and components industries (State Street). This feature allows investors to take strategic or tactical positions at a more targeted level than traditional style based investing. The companies in the XLK are selected from the S&P 500 based on the Global Industry Classification Standard (GICS) where companies are assigned to a single classification based on revenue.

2.2 Data Collection

Using the dataset from the investing website, we derive the following variables for the analysis:

- **Price** (in USD): The final price at which the ETF traded when the stock exchange closed (typically 4:00 PM EST).
- **Open** (in USD): The price of the very first transaction when the market opened (9:30 AM EST).
- **High & Low** (in USD): The maximum and minimum prices reached during the trading session.
- **Vol.:** The total number of shares traded during the day (e.g., 15.90M means 15.9 million shares).
- **Change:** The percentage difference between the current day’s closing price and the previous day’s closing price.

The dataset comprises daily observations spanning from January 3, 2017, to March 28, 2026. This time frame was specifically selected to ensure that the launch of ChatGPT (November 30, 2022) serves as a central pivot point, providing a balanced sample size for comparing pre- and post-innovation market dynamics

2.3 Data Cleaning

In the raw dataset from the investing website, we observe the point information on weekends are carried over from the Friday of the same week with empty number in **Vol.**. We removed these points since they are not consistent with the general pattern and can be misleading in identifying the long-term price fluctuation. We then splitted the dataset into two groups: **before_data** and **after_data**, using the time stamp November 30, 2022 (the release date of chatGPT) as the cutoff.

Table 1: Summary Statistics of Stock Price

mean_price	sd_price	min_price	max_price
69.91683	33.95005	24.4	152.07

Table 1 summarized the descriptive statistics for the stock price from January 3, 2017, to March 27, 2026, including the mean, standard deviation, and range. A more intuitive historical price trajectory is illustrated in Figure 1, where the vertical red dashed line marks the official release




Figure 1: Raw Stock Price Over Time




Figure 2: Autocorrelation Function (ACF) of Raw Stock Price.

of ChatGPT. Overall, the stock shows a sustained upward trend throughout the observed period.

The autocorrelation function (ACF) plot in Figure 2 demonstrates a strong correlation between data points, indicating this time series is non-stationary and does not exhibit white noise characteristics. The ACF values persist well beyond the significance boundaries and decay at a very slow rate. This high correlation is explainable since the stock prices are accumulated based on the previous day's close price. Therefore, data transformation is required to achieve stationarity prior to modeling.




Figure 3: Daily Log Returns of Stock Price (2017–2026).

To stabilize the variance and ensure stationarity, we applied a natural logarithmic transformation followed by first-order differencing. Figure 3 shows the log return over time after transformation. While the log return mostly fluctuates around a mean of zero with small fluctuations, there are few large spikes. The biggest fluctuations occur around 2020 when the pandemic began, suggesting that the pandemic can cause substantial uncertainty and sharp movements in the market. Although Figure 4 shows there still exist a few autocorrelation coefficients outside the significance boundaries, the Augmented Dickey-Fuller (ADF) test was conducted for a formal assessment. The test yielded a p-value of 0.01, providing strong empirical evidence to reject the null hypothesis of a unit root and confirming that the log return series is stationary.




Figure 4: Autocorrelation Function (ACF) of Log Return.

3 Method

3.1 Compare log return using regression

Table 2: Summary Statistics of Stock Price

period	Mean	SD
After AI	-0.00078	0.01456
Before AI	-0.00069	0.01641

According to table Table 2, the mean log return before and after the appearance of ChatGPT is both close to zero. To test whether there is a difference on log return before and after ChatGPT appearance, a regression model with a dummy variable D_t as the predictor was fitted.

$$\log(\text{return}) = \beta_0 + \beta_1 D_t$$

where $D_t = 0$ is before the appearance of ChatGPT, and $D_t = 1$ is after the appearance of ChatGPT. A two sample mean test can be used to determine whether D_t is significant. The null hypothesis is that the mean log returns in the two periods are equal, while the alternative hypothesis is that they are different.


Table 3: Model Coefficients



	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.001	0.000	-1.684	0.092
AI_dummy	0.000	0.001	-0.132	0.895



The p value of the dummy variable is 0.895, which is greater than 0.05. Thus, we fail to reject the null hypothesis, and conclude that there is no difference on the mean log return values before and after the appearance of AI.

3.2 Fit ARIMA Model

Using the ACF and PACF plots, we can determine the appropriate order of the ARIMA model. According to Figure 4, the ACF gradually tails off, suggesting the presence of autoregressive behavior. According to `fig-logpacf`, there are strong spikes at lags 2, 4, 7, and 8, and several other significant lags. Based on these patterns, ARIMA(1,1,1), ARIMA(2,1,2), ARIMA(4,1,4), ARIMA(7,1,7), and ARIMA(8,1,8) models were fitted to the price data to capture the time series.







Model Diagnosis

The p values for Ljung-Box Statistic plots indicate that ARIMA(1,1,1), ARIMA(2,1,2), and ARIMA (4,1,4) are not good fits, because p values are significant for all lags. ACF plots for ARIMA(7,1,7) and ARIMA (8,1,8) looks stationary, but QQ plots show the points are deviated from the reference line, meaning the normality condition is not met. To determine which model performs the best, we compare the model AIC, AICc, and BIC values.

Model Choice

Table 4: Comparison of Model Performance

	AIC	AICc	BIC
ARIMA(1,1,1)	-5.477658	-5.477656	-5.470223
ARIMA(2,1,2)	-5.502680	-5.502673	-5.490290
ARIMA(4,1,4)	-5.503696	-5.503669	-5.481393
ARIMA(7,1,7)	-5.503852	-5.503774	-5.466680
ARIMA(8,1,8)	-5.501314	-5.501212	-5.459185

The ARIMA(7,1,7) has the smallest AICc value among all the models. Therefore, we pick ARIMA(7,1,7) as our best model.

Forecasting

We are interested in how significant the data after 2022/11/30, when ChatGPT got launched, is to forecasting future stock market prices. We will be using data frames of different time sets: 2017/1/3 - 2026/3/27, 2017/1/3 - 2022/11/30 and 2022/11/30 - 2026/3/27. (Date) The forecast models we will be using three of the “benchmark” methods: Mean method, Naive method and Seasonal naive method. (describe methods) Drift method and decomposition method seem unnecessary with log return data because we have already removed the trend.

First, we aimed to examine whether AI has had an unexpected impact on the stock market and contributed to its fluctuations. We would like to use data from 2017/1/3 - 2022/11/30 to forecast data after 2022/11/30.

Table 5: Forecast Errors

Table 5: Forecast Errors

Method	MAE	RMSE
Mean	0.01049	0.01457
Naive	0.01241	0.01597
Seasonal	0.02161	0.02880



According to Table 5(explain meaning), all the scaled dependent errors are small, even for these easier approaches. It appears that the data prior to 2022/11/30 can effectively predict the data following the cutoff point. Furthermore, this suggests that AI have not yet contribute a significant impact in the stock market. As a result, we will use the entire time frame, allocating 80% for training and 20% for testing, to build a forecast model in predict prices in 2026.

Table 6: Forecast Errors-full model

Table 6: Forecast Errors-full model


Method	MAE	RMSE
Mean	0.01128	0.01631
Naive	0.01459	0.01845
Seasonal	0.01311	0.01754

Table 6 shows the mean method is the best with lowest MAE and RMSE. We did not put MAPE here for comparison, because our log return data moves around 0, where some small values here make MAPE enormous and not as a good parameter to look at.



Next, we would like to also make a prediction interval for this forecast model.[ACF:()]. The histogram looks normal, however qq-plot reveals that log return does not satisfy the linearity assumption. This suggests that our prediction interval is off.

(Will do ARIMA but also keeping this)



Summary

This project tests whether the rise of AI has significantly affected the behavior of technology sector stock returns. Using the XLK ETF price data from 2019 to 2026, we transformed the series into log returns to make the data stationary and applied AR(1), MA(1), and ARMA (1,1) time series models. The result suggests that stock return pattern remained relatively stable before and after the emergence of AI, indicating that AI has not yet had a significant impact on the technology stock market price. More broadly, this shows that while AI is often seen as a transformative force, its impact on financial markets may not be immediately reflected in stock return fluctuations.

Limitations

There are some potential limitations for the dataset. One of the limitations is that there's possibility for State Street to rebalance the stock composition over time, which means the dataset may not remain fully consistent throughout the study period. The companies in the XLK ETF are selected based on the GICS. However, the GICS structure is reviewed annually, meaning State Street need to adjust the XLK portfolio at least on a yearly base. This could reduce the reliability of comparing stock performance across different years and weaken the robustness of analysis.

Another limitation is the dependency of the dataset. In the data analysis, we noticed few large spikes in the log return plot, suggesting factors like pandemics may also attributes to the stock price fluctuation. In Section 3, we found out the time series data points remain highly correlated even after the data transformation. For instance, the QQ plot indicates deviations from normality, suggesting that model assumptions may not hold and prediction intervals may be unreliable. In addition, our model focuses on the short-term impacts after the release of ChatGPT, but the effect of AI may be better reflected in a long term trend rather than a boom. The log transformation may mask the underlying long term trend, because it removes the trend to make the data stationary. The log transformation only focuses on daily changes rather than long-term structural changes. As a results, model based on the log transformation is better at studying short-term market reaction but may be less effective at capturing long-term changes in the market over time.

Also, AR, MA, ARMA are all models that are only a function of the time series itself. They do not include other important factors, such as government policy or economic policy changes. There are many other factors that affects the stock market that we did not incorporate, such as economic indicators and interest rates. Since stock market is influenced by many external variables, excluding these factors may limit the strength and forecasting accuracy of the models.

Future Directions

To account for external factors that might affect the stock market, we could add these as predictors in our model. This can help us better estimate stock prices.

Although we have not been able to conclude AI has significant impact on the behavior of technology sector stock returns, it is still useful to estimate how sensitive technology sector stock price comparing to the broader market indices or other sectors. We might expect the technology sector to show higher volatility and respond more to growth expectations. In future studies, we can fit and compare separate models for different market segments, such as the S&P 500 and individual sectors like healthcare or financials, to better understand these differences.

It is also possible to incorporate the method proposed by Campos-Martins (2024) with our model. After identifying notable development by AI companies, we can adjust our predictors to improve accuracy. For example, we could include an indicator variable that captures major AI events or achievements. This would allow the model to account for changes over time driven by new innovations, instead of relying only on fixed predictors.

Finally, our research should follow a consistent process where we constantly refit the models as the stock market changes over time. Updating the model with new data can help us see whether there are any long-term effects or clear shifts in the data, especially as AI continues to develop and become more widely used.